

QProxy

Towards running legacy applications inside AWS Nitro Type Enclaves

Jalil David Salamé Messina

Advisor: Dr. Quoc Do Le

Chair of Computer Systems


<https://dse.in.tum.de/>



15.01.2024 – 15.04.2024

Motivation: Research context

- Confidential applications need isolation
- Shared cloud environments don't ensure it
- Trusted Execution Environments (TEEs) do
 - Application Level
 - Intel SGX
 - Virtual Machine (VM) Level:
 - Intel TDX
 - AMD SEV-SNP
 - AWS Nitro enclaves
 - Huawei Qingtian enclaves
- AWS Nitro and Huawei Qingtian enclaves only allow VSocket connections
- We need a TCP ↔ VSocket translation layer



- Nitriding Proxy¹
- Enclaver (<https://github.com/edgebitio/enclaver>)
- Fortanix (Confidential Computing Manager)
(<https://support.fortanix.com/hc/en-us/categories/360003107511-Confidential-Computing-Manager>)

¹ Winter, Philipp, et al. "A Framework for Building Secure, Scalable, Networked Enclaves." arXiv preprint arXiv:2206.04123 (2022).

- Limitations of State-of-the-art:
 - Focused only on AWS enclaves
 - Integrates with AWS KMS only
 - Automatic TLS Termination
 - Poor Performance
- Our work:
 - Generic Proxy
 - No direct integrations
 - No TLS termination
 - Acceptable performance

- In this Thesis, our goal is to develop and implement a proxy system that:
 - Exposes enclave services outside the enclave
 - Supports both AWS Nitro and Huawei Qingtian enclaves
 - Without modifying the confidential application
 - With minimal overhead
 - Scale to high loads

QProxy: A TCP to VSocket proxy for enclaves



QProxy

QProxy design goals:

- Security
- Transparency
- Performance
- Scalability

- ~~Motivation~~

- Background

- Trusted Execution Environments (TEEs)
- Nitro-type Enclaves

- Design

- Implementation

- Evaluation

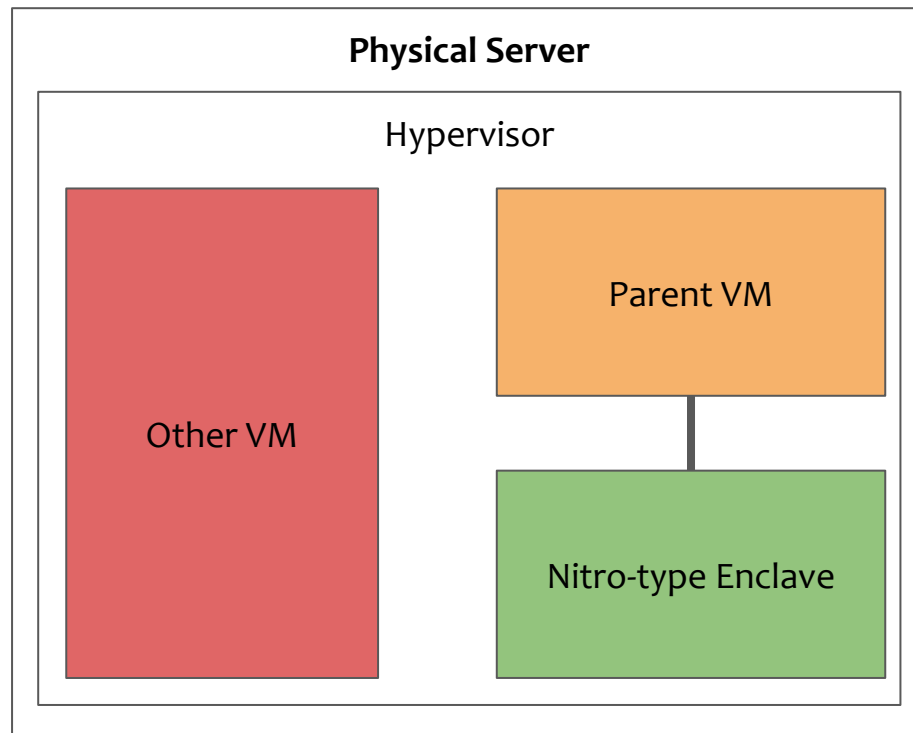
Trusted Execution Environments (TEEs)



- AMD SEV-SNP/Intel TDX
 - VM Level
 - Confidential (through memory encryption)
 - Integrity (through memory protections)
 - Remote Attestation
- Intel SGX
 - Application Level
- ARM TrustZone
 - VM Level (boot time)
- ARM CCA
 - VM Level (runtime)

Nitro-type Enclaves (AWS Nitro and Huawei Qingtian)

- VM as a TEE
- Isolated (besides VSocket)
 - Mitigates Spectre & Meltdown
- Signed Attestation Document



Outline

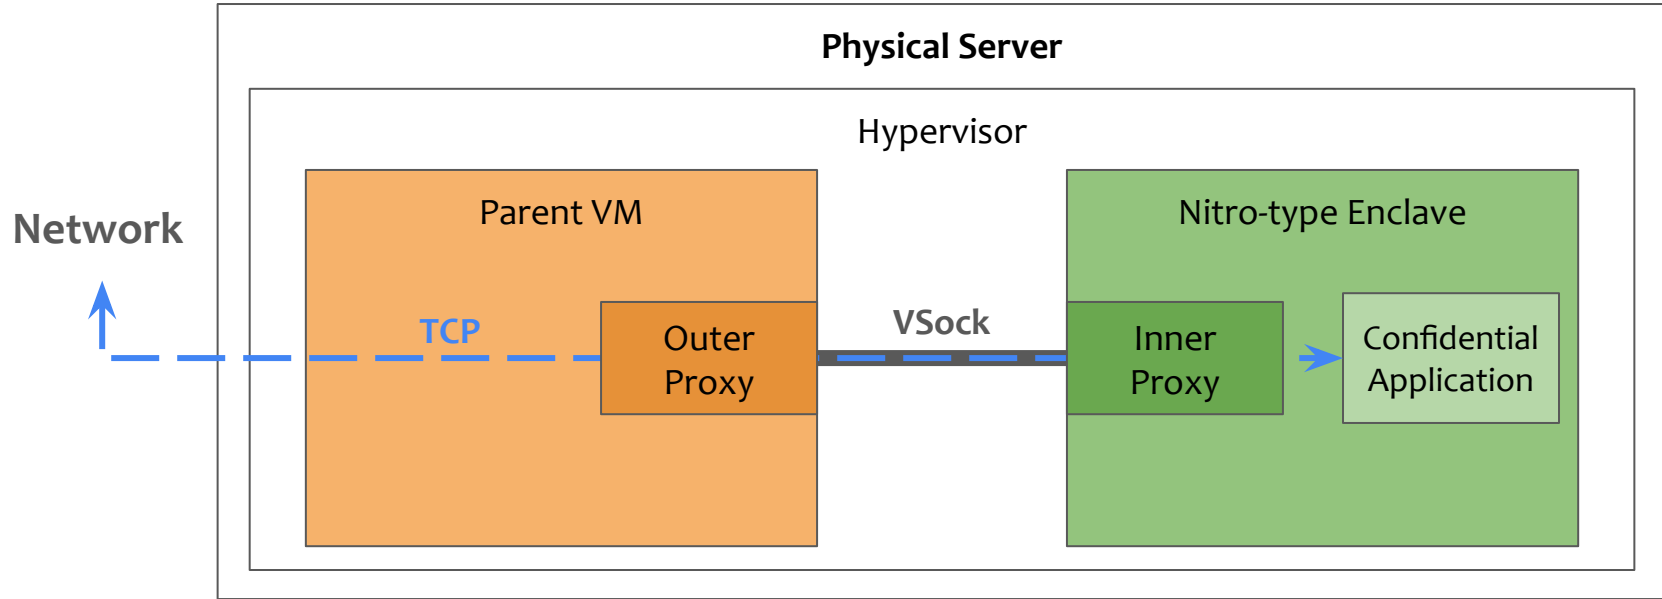
- ~~Motivation~~
- ~~Background~~
- Design
 - System design
- Implementation
- Evaluation

System overview



- Outside (host) Proxy
 - Network to and from Enclave
 - Rudimentary access control
- Inside (enclave) Proxy
 - Enclave to and from applications(s)

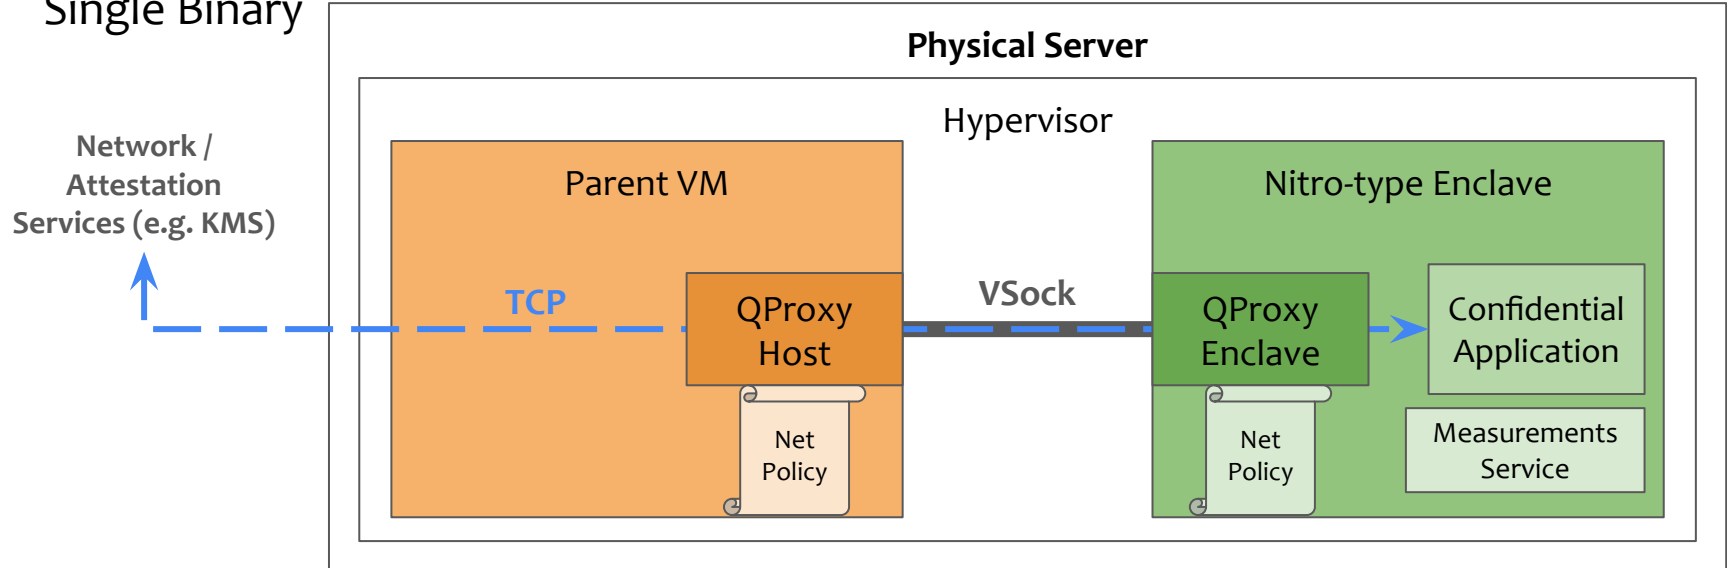
Design overview



QProxy transparently establishes connections to the confidential application

Design details

- Layer 4 Proxy (Over TCP)
- No Tunnelling
- Rust with Tokio.rs
- Single Binary



Outline

- ~~Motivation~~
- ~~Background~~
- ~~Design~~
- Implementation
- Evaluation

QProxy is built with **Rust**¹ and **Tokio.rs**²

- Memory-safety at compile-time
- No unsafe Rust
- Leveraging well-tested OSS libraries
- Highly concurrent through Tokio.rs
- Currently running in production

¹Rust Programming Language: <https://www.rust-lang.org/>

²Tokio an asynchronous runtime for Rust: <https://tokio.rs/>

Outline

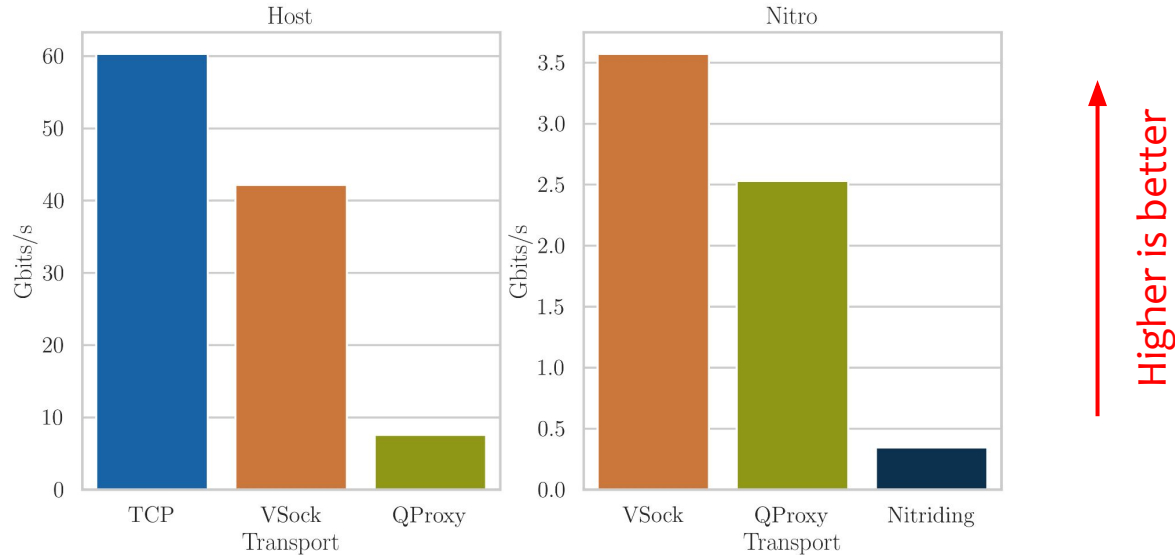
- ~~Motivation~~
- ~~Background~~
- ~~Design~~
- ~~Implementation~~
- Evaluation

- What is the performance overhead of QProxy?
 - Latency
 - Throughput
- Is it transparent?
 - Run Redis through QProxy
- How does it compare with the state-of-the-art?
 - Compare against Nitriding
- How does it scale with the available resources?
 - Increasing the available resources

- Experimental setup (Nitro):
 - c6.8xlarge AWS Nitro EC2
 - Intel® Xeon® Platinum 8375C CPU @ 2.90GHz
 - 32 vCPUs (2 allocated to the enclave)
 - 64 GiB DRAM (1.5 GiB allocated to the enclave)
- Experimental setup (Qingtian):
 - c7.8xlarge.4 Huawei EC2
 - Intel® Xeon® Platinum 8378A CPU @ 3.00GHz
 - 32 vCPUs (2 allocated to the enclave)
 - 128 GiB DRAM (1 GiB allocated to the enclave)
- Baseline:
 - Using a VSocket service (no proxy overhead)
 - Using Nitriding (state-of-the-art) (only on Nitro)

Performance overhead (throughput)


iperf3-vsock: single-core, throughput stress-test (5 minute average)



QProxy is between 2.5 and 5 times better than Nitriding

Performance overhead (throughput)


memtier_benchmark: redis load-test (average over 5 minutes)



QProxy achieves slightly better throughput than Nitriding

Performance overhead (latency)

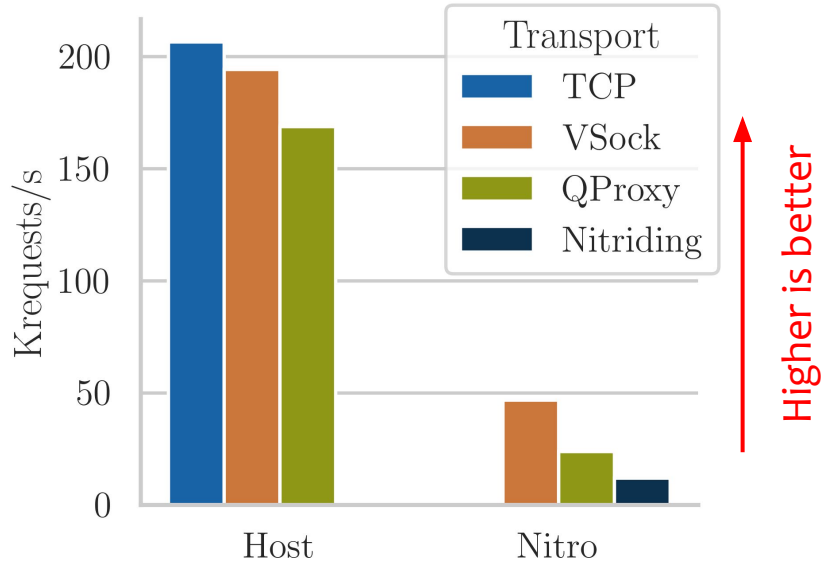
memtier_benchmark: redis load-test, rate-limited to 10kops/s (5 minutes)



QProxy achieves circa 1.5 times better latency than Nitriding

Performance overhead (throughput)

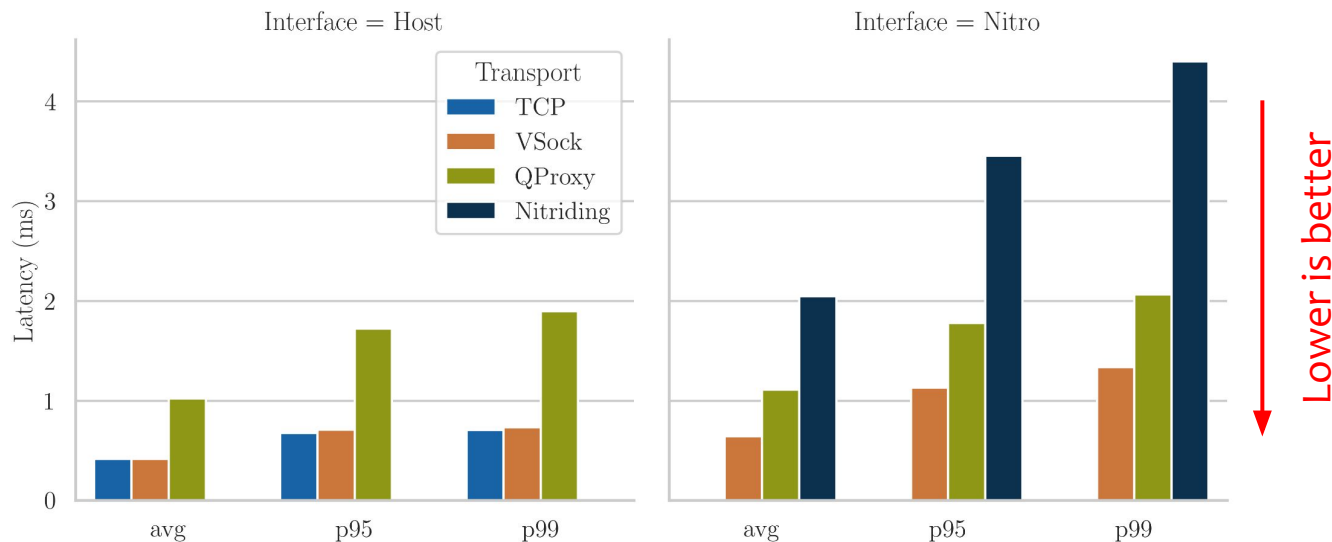
oha: HTTP load-test, 50 connections over 30 threads (5 minute average)



QProxy achieves circa 2 times the throughput of Nitriding

Performance overhead (latency)


oha: HTTP load-test, 50 connections over 30 threads, limited to 10krps (5 minutes)



QProxy achieves circa 2 times better latency than Nitriding

Scaling (throughput)


iperf3-vsock: single-core, throughput stress-test (5 minute average)



AWS Nitro Enclaves have decreased baseline performance with more cores

Scaling (throughput)

oha: HTTP load-test, 50 connections over 30 threads (5 minute average)



This problem is also present in Qingtian enclaves

Current applications are not designed for Nitro-type enclaves

- VSocket transport
- confidentiality & integrity
- remote attestation

QProxy:

- transparent VSocket proxy
- no source code modifications
- Circa 2x performance of Nitriding

Future Work:

- Integrate with configs-tsm (kernel remote attestation ABI)
- Integrate with KMS services

Questions?