

Focaccia

Execution State Comparison for Emulators using Symbolic Execution

Nicola Crivellin

Advisor: Sebastian Reimers & Theofilos Augoustis

Chair of Computer Systems

<https://dse.in.tum.de/>

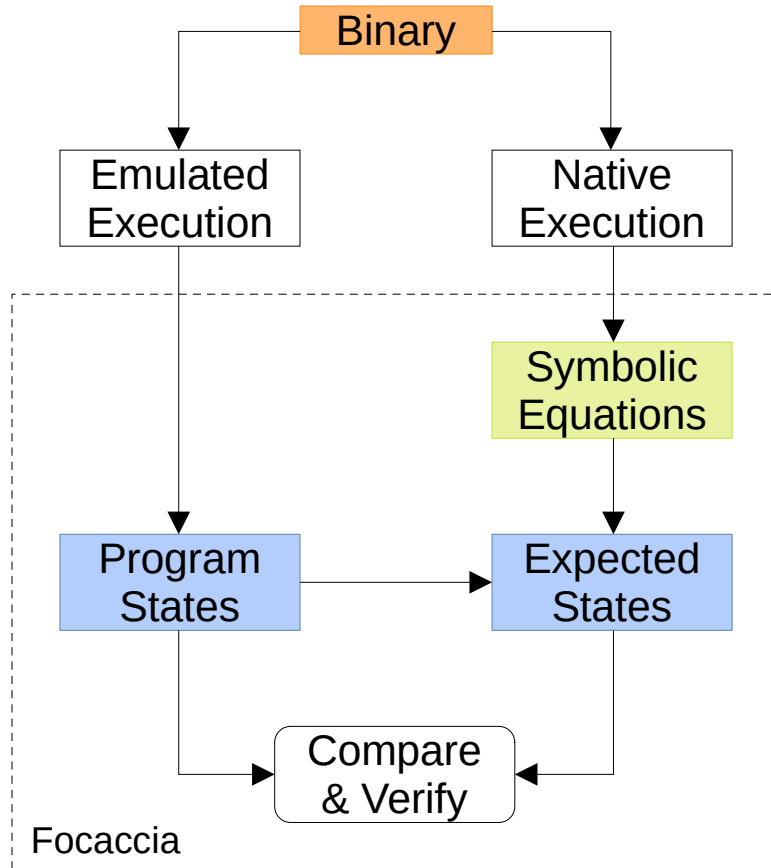


15.11.2023 – 15.05.2024

- Developing emulators is difficult
- We want to make it easier to test them
 - Systematic
 - Automatic
 - Simple
- Premise: Emulators output code

Can we use symbolic execution to verify an emulator's output?

The system: Overview



- Automatic testing system for emulators (currently for x86 guests)
- Calculates test truths via oracle
 - Uses symbolic execution
- Compares truth states to emulated states
 - Figures out whether instructions modify state correctly

The system: Recording a concrete trace

Emulated execution

0x401032
0x401035
0x401038
0x40103f
0x401043
...

Snapshot at instruction 0x401038

Registers

RAX: 0x0	RBX: 0x0
RCX: 0x0	RDY: 0x0
RSI: 0x0	RDI: 0x7fffffffef680
RBP: 0x0	RSP: 0x7fffffffef680
...	

Memory

The system: Generating a symbolic trace

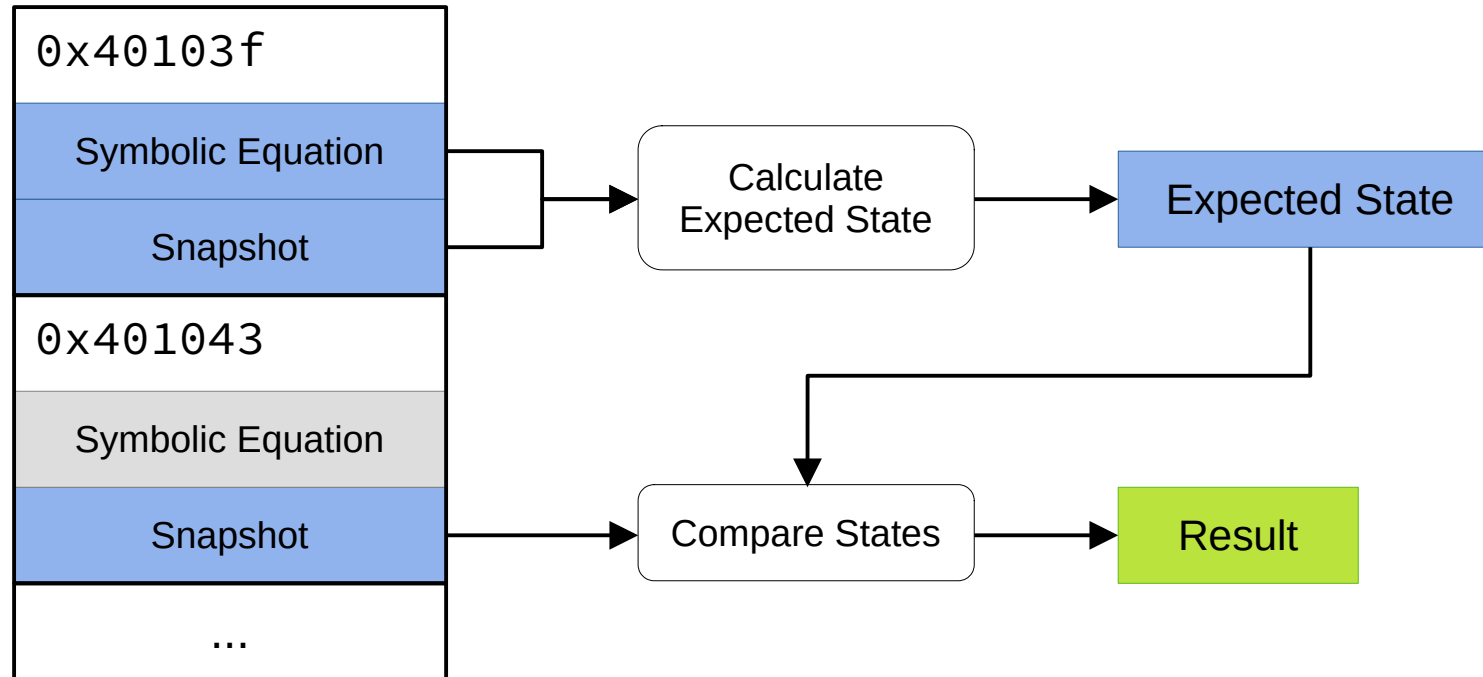
Native execution

0x401032
0x401035
0x401038
0x40103f
0x401043
...

Symbolic equation for AND RSP, 0xFFFFFFFFFFFFFFF0

zf = (RSP & 0xFFFFFFFFFFFFFFF0) ? (0x0, 0x1)
nf = (RSP & 0xFFFFFFFFFFFFFFF0) [63:64]
pf = parity(RSP & 0xF0)
of = 0x0
cf = 0x0
RSP = RSP & 0xFFFFFFFFFFFFFFF0
RIP = 0x401043

The system: Predicting outcomes



Example

```
$ qemu-x86_64 -g 12345 bug.out &  
$ ./capture_transforms.py -o symbols.trace bug.out  
$ ./verify_qemu.py --symb-trace symbols.trace localhost 12345
```

```
[...]
```

```
-----  
For PC=0x40116a  
-----
```

1. [ERROR] Content of register RAX is false. Expected value: 0x1234567812345678, actual value in the translation: 0x12345678.

Expected transformation: Symbolic state transformation 0x40116a -> 0x40116e:

[Symbols]

ZF = (~@32[RBP + 0xFFFFFFFFFFFFFFEC]) == (~RAX[0:32])

RAX = ((~@32[RBP + 0xFFFFFFFFFFFFFFEC]) == (~RAX[0:32]))?(RAX,{@32[RBP + 0xFFFFFFFFFFFFFFEC] 0 32, 0x0 32 64})

RIP = 0x40116E

@32[RBP + 0xFFFFFFFFFFFFFFEC] = ((~@32[RBP + 0xFFFFFFFFFFFFFFEC]) == (~RAX[0:32]))?(RDI[0:32],@32[RBP + 0xFFFFFFFFFFFFFFEC])

[Instructions]

CMPXCHG DWORD PTR [RBP + 0xFFFFFFFFFFFFFFEC], EDI

Actual difference: Snapshot (x86_64): {'RIP': '0x4', 'RBP': '0x0', 'RAX': '-0x1234567800000000'}

Example

```
$ qemu-x86_64 -g 12345 bug.out &  
$ ./capture_transforms.py -o symbols.trace bug.out  
$ ./verify_qemu.py --symb-trace symbols.trace localhost 12345
```

[...]

For PC=0x40116a

1. [ERROR] Content of register RAX is false. Expected value: 0x1234567812345678, actual value in the translation: 0x12345678.

Expected transformation: Symbolic state transformation 0x40116a -> 0x40116e:

[Symbols]

ZF = (~@32[RBP + 0xFFFFFFFFFFFFFFFF]) == (~RAX[0:32])

RAX = ((~@32[RBP + 0xFFFFFFFFFFFFFFFF]) == (~RAX[0:32]))?(RAX,{@32[RBP + 0xFFFFFFFFFFFFFFFF] 0 32, 0x0 32 64})

RIP = 0x40116E

@32[RBP + 0xFFFFFFFFFFFFFFFF] = ((~@32[RBP + 0xFFFFFFFFFFFFFFFF]) == (~RAX[0:32]))?(RDI[0:32],@32[RBP + 0xFFFFFFFFFFFFFFFF])

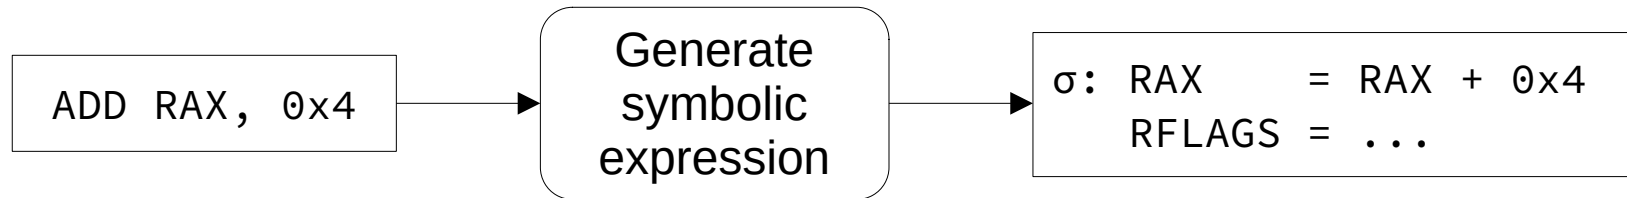
[Instructions]

CMPXCHG DWORD PTR [RBP + 0xFFFFFFFFFFFFFFFF], EDI

Actual difference: Snapshot (x86_64): {'RIP': '0x4', 'RBP': '0x0', 'RAX': '-0x1234567800000000'}

- Search QEMU's Gitlab for instruction-semantics bugs with x86 guest code
- Data: 9 bugs since 2021
- Procedure: Reproduce bugs, then run them through Focaccia
- Result: 1/9 bugs found successfully
 - Found: 1 bug
 - Symbolic execution backend: 6 bugs
 - QEMU crashes: 1 bug
 - Not reproducible: 1 bug

- Symbolic execution backend acts as a trusted emulator
- Focaccia uses one correct implementation to verify all other emulators
 - Software oracle vs. hardware oracle?
- Approach is not systematic: works on specific test cases



- Gives useful results if the backend is good
- Resilient to faults: backend, input data, emulator
- Tool is highly usable: minimally invasive, out-of-the-box paradigm, configurable